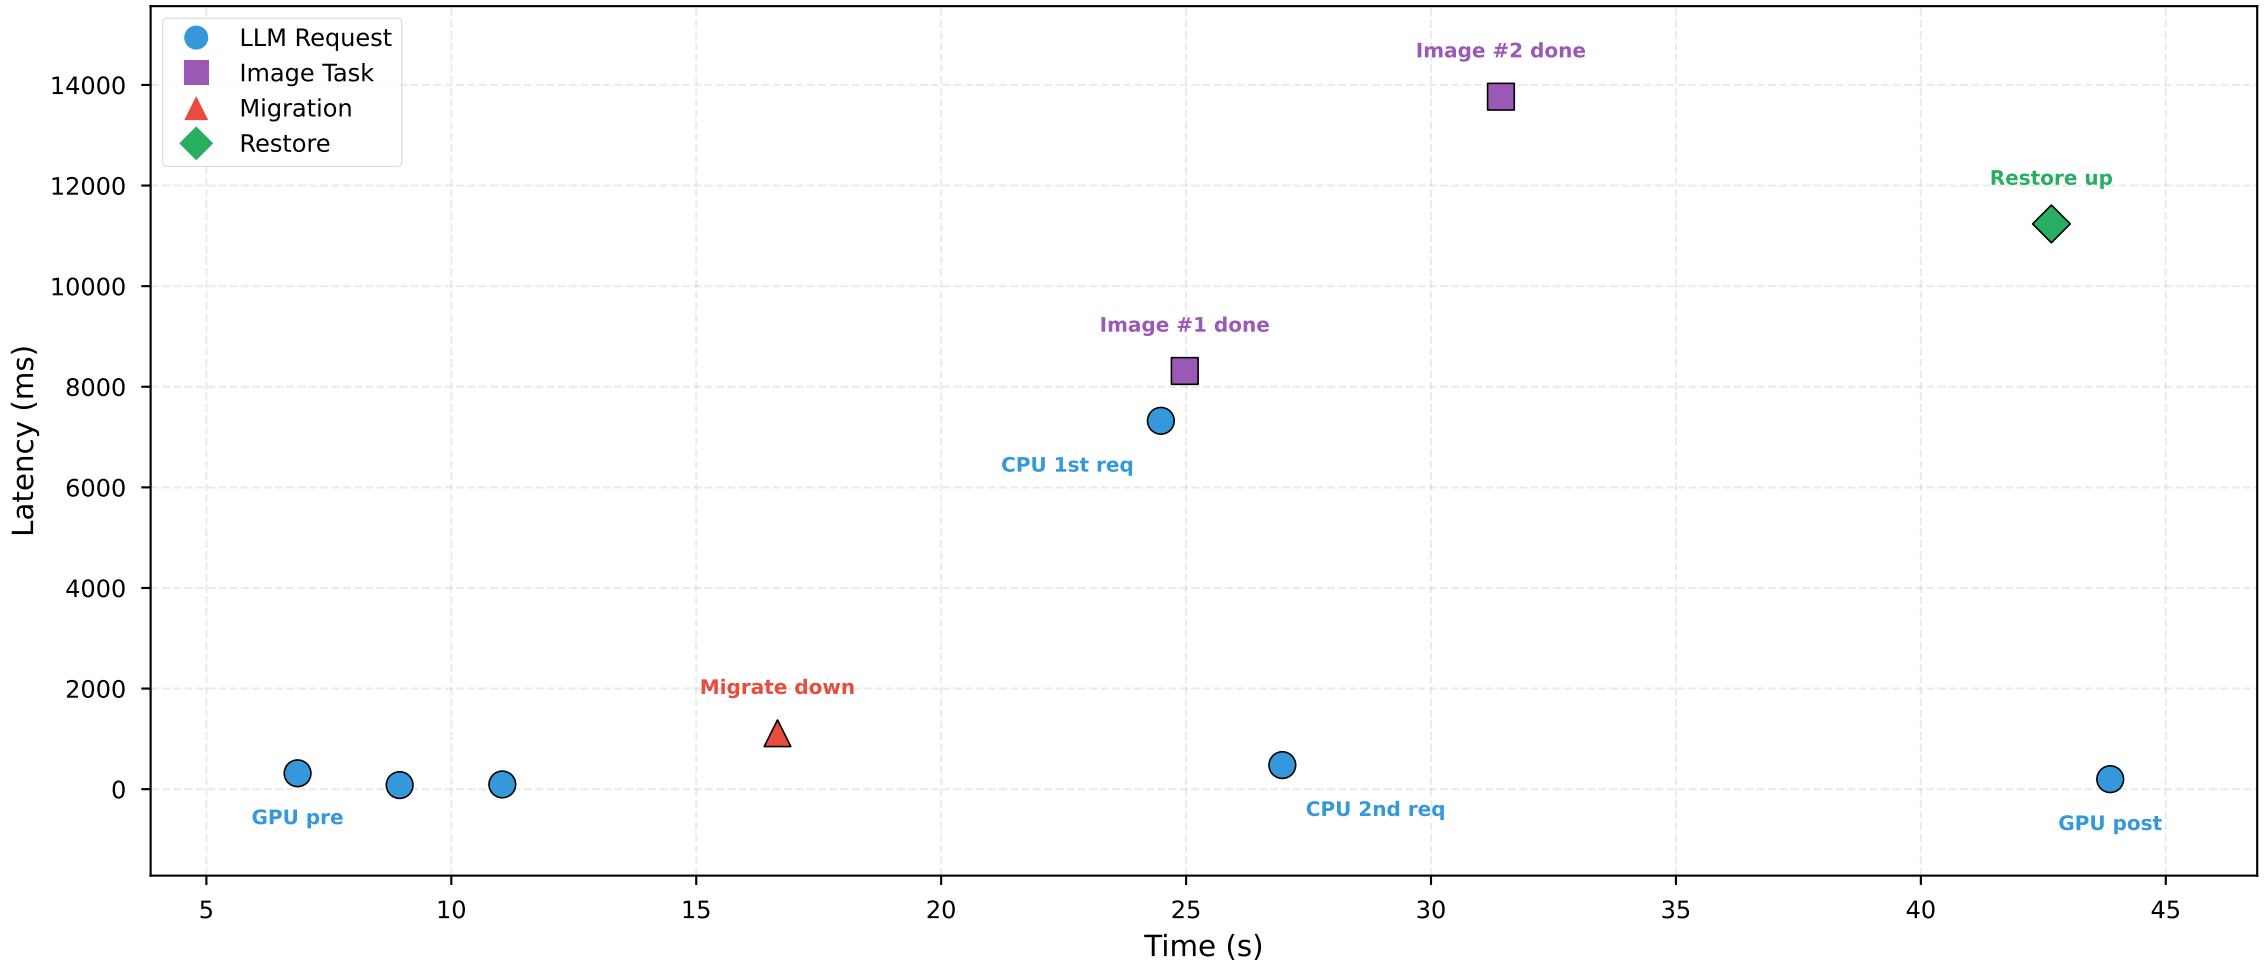
Real Workload Trace: LLM Chat + Image Generation
LLM Request
Image Task
Migration
Restore
Image #2 done
Restore up
Image #1 done
CPU 1st req
Migrate down
GPU pre
CPU 2nd req
GPU post
Latency (ms)
Time (s)
0
2000
4000
6000
8000
10000
12000
14000
5
10
15
20
25
30
35
40
45